

# Tutorial 1

Python  
Tokenisation and Segmentation  
Indexing and Search  
Evaluation

# Python

- Many ways to learn – here are some resources:
  - <https://docs.python.org/3/tutorial/>
  - <http://www.slideshare.net/MattHarrison4/learn-90>
- We'll visit a Sheffield tutorial on this

# Tokenisation

- Goal: convert stream of bytes to tokens
- Simple method: Python's `.split()` function
- Let's get some data and play with it!
  - `import nltk`
  - `nltk.download()`
- Go to *Corpora*, find the *Gutenberg* corpus
- Let's load it
  - `from nltk.corpus import gutenberg`
  - `text = gutenberg.raw("carroll-alice.txt")`
  - `print("\n".join(gutenberg.fileids()))`

# Whitespace Tokenisation

- Assume that tokens are whitespace-delimited
- Implement a tokeniser using `split()`
- Write a program that:
  - Loads and tokenises the text
  - Prints the number of tokens it contains

# Word Tokenisation

- Now use NLTK's `word_tokenize()` function to find token boundaries
  - `nltk.download()`
  - Get `punkt` – a tokenisation package (for words and sentences)
- Write a program
  - That reads and tokenises the text
  - Prints the number of tokens it contains.
- How does this compare to using whitespace tokenisation? Why?

# Chinese Word Segmentation

- Chinese has no explicit word boundaries

中文句子由连续的一系列单词组成

- Greedy Matching
  - Use a list of known words
  - Find longest possible matches

# Chinese Word Segmentation

- Algorithm:
  - Start at the beginning of the sentence
  - Find longest sequence of (up to  $n$ ) consecutive characters in the list
  - If a match is found, assume that's the next word
    - Store it
    - Move on
  - Otherwise, assume a single-character word
  - Continue until the end of the sentence

# Chinese Word Segmentation

- Download
  - Unsegmented text
  - Word list
- Evaluate performance using the evaluation script
  - What effect does tuning  $n$  have?



# Indexing

- Goal is to build a word-to-document reference
- We can store this with shelve
  - `import sys`
  - `index_filename = sys.argv[1]`
  - `import shelve`
  - `with shelve.open(index_filename, 'c') as index:`
    - `index['x'] = [1,2,3]`
    - `index.sync()`
- This means we only need to do our indexing once

# Indexing

- Write a program that:
  - Takes an index filename
  - Opens the gutenber corpus
  - Goes through the documents in the corpus, tokenising each one
    - `for document_id in gutenber.fileids()`
      - `tokens = nltk.word_tokenize(gutenber.raw(document_id))`
  - For each token, adds an entry in the index containing that document's ID
    - `if term not in index:`
      - `Index[term] = []`
    - `index[term].append(document_id)`

# Searching

- Open the index
- Take a query string
  - `sys.stdin.readline()`
  - `.strip()`
- Tokenise the query string
- Find matching documents
  - Look up the entries in the shelved index
    - Lists are not unique
    - `set(list)`

# Ranking

- Update the indexing program to include TF
  - `term_count = document_tokens.count(term)`
  - `index[term].append( (document_id, term_count) )`
- DF is just `len(index[term])`
- $TF.IDF = TF / DF$ 
  - Variants:
    - +1 smoothing to IDF
    - Take logs on both sides
- Calculate TFIDF for all documents in list, then rank
- Congratulations! Your own pre-Google search engine

# Evaluation

- Evaluating IR:
  - How much did you find?
  - How much did you find, that shouldn't have been there?
  - How much did you miss?
- Bonus exercise:
  - Repeat the above, for Chinese or Russian!
  - Xinhua, Russia Today

# Assignment

- Read in a collection; run some queries over it
- Total 500~1500 words
  - Discuss performance
  - Identify problems
  - Suggest methods for improvement
- Include a copy of your source code
  - For an A grade, extend the system beyond this specification, and describe your extension
  - Due October 6
  - Mail me: **leonderczynski@gmail.com**